

Tutorial 1

This tutorial is an introduction to backpropagation and TensorFlow. It starts by explaining the basics of backpropagation, including how to calculate gradients and update weights. It then shows how to implement backpropagation in TensorFlow. Overall, the tutorial is intended to give readers the tools they need to start building their own neural networks.

Tutorial 2

This tutorial is about using transfer learning with TensorFlow for object classification. It covers the basics of transfer learning and how it can be used to improve the performance of a neural network. Additionally, the tutorial discusses some best practices, such as choosing the right pre-trained model and fine-tuning certain layers of the model.

Tutorial 3

The third tutorial explains the basics of semantic segmentation and how the U-Net architecture is designed to perform this task. Also, the tutorial discusses some practices for training a U-Net, such as using data augmentation and using a loss function that accounts for class imbalance. Overall, the tutorial shows how to use a U-Net to perform this task in TensorFlow.